

Cassandra and Kafka

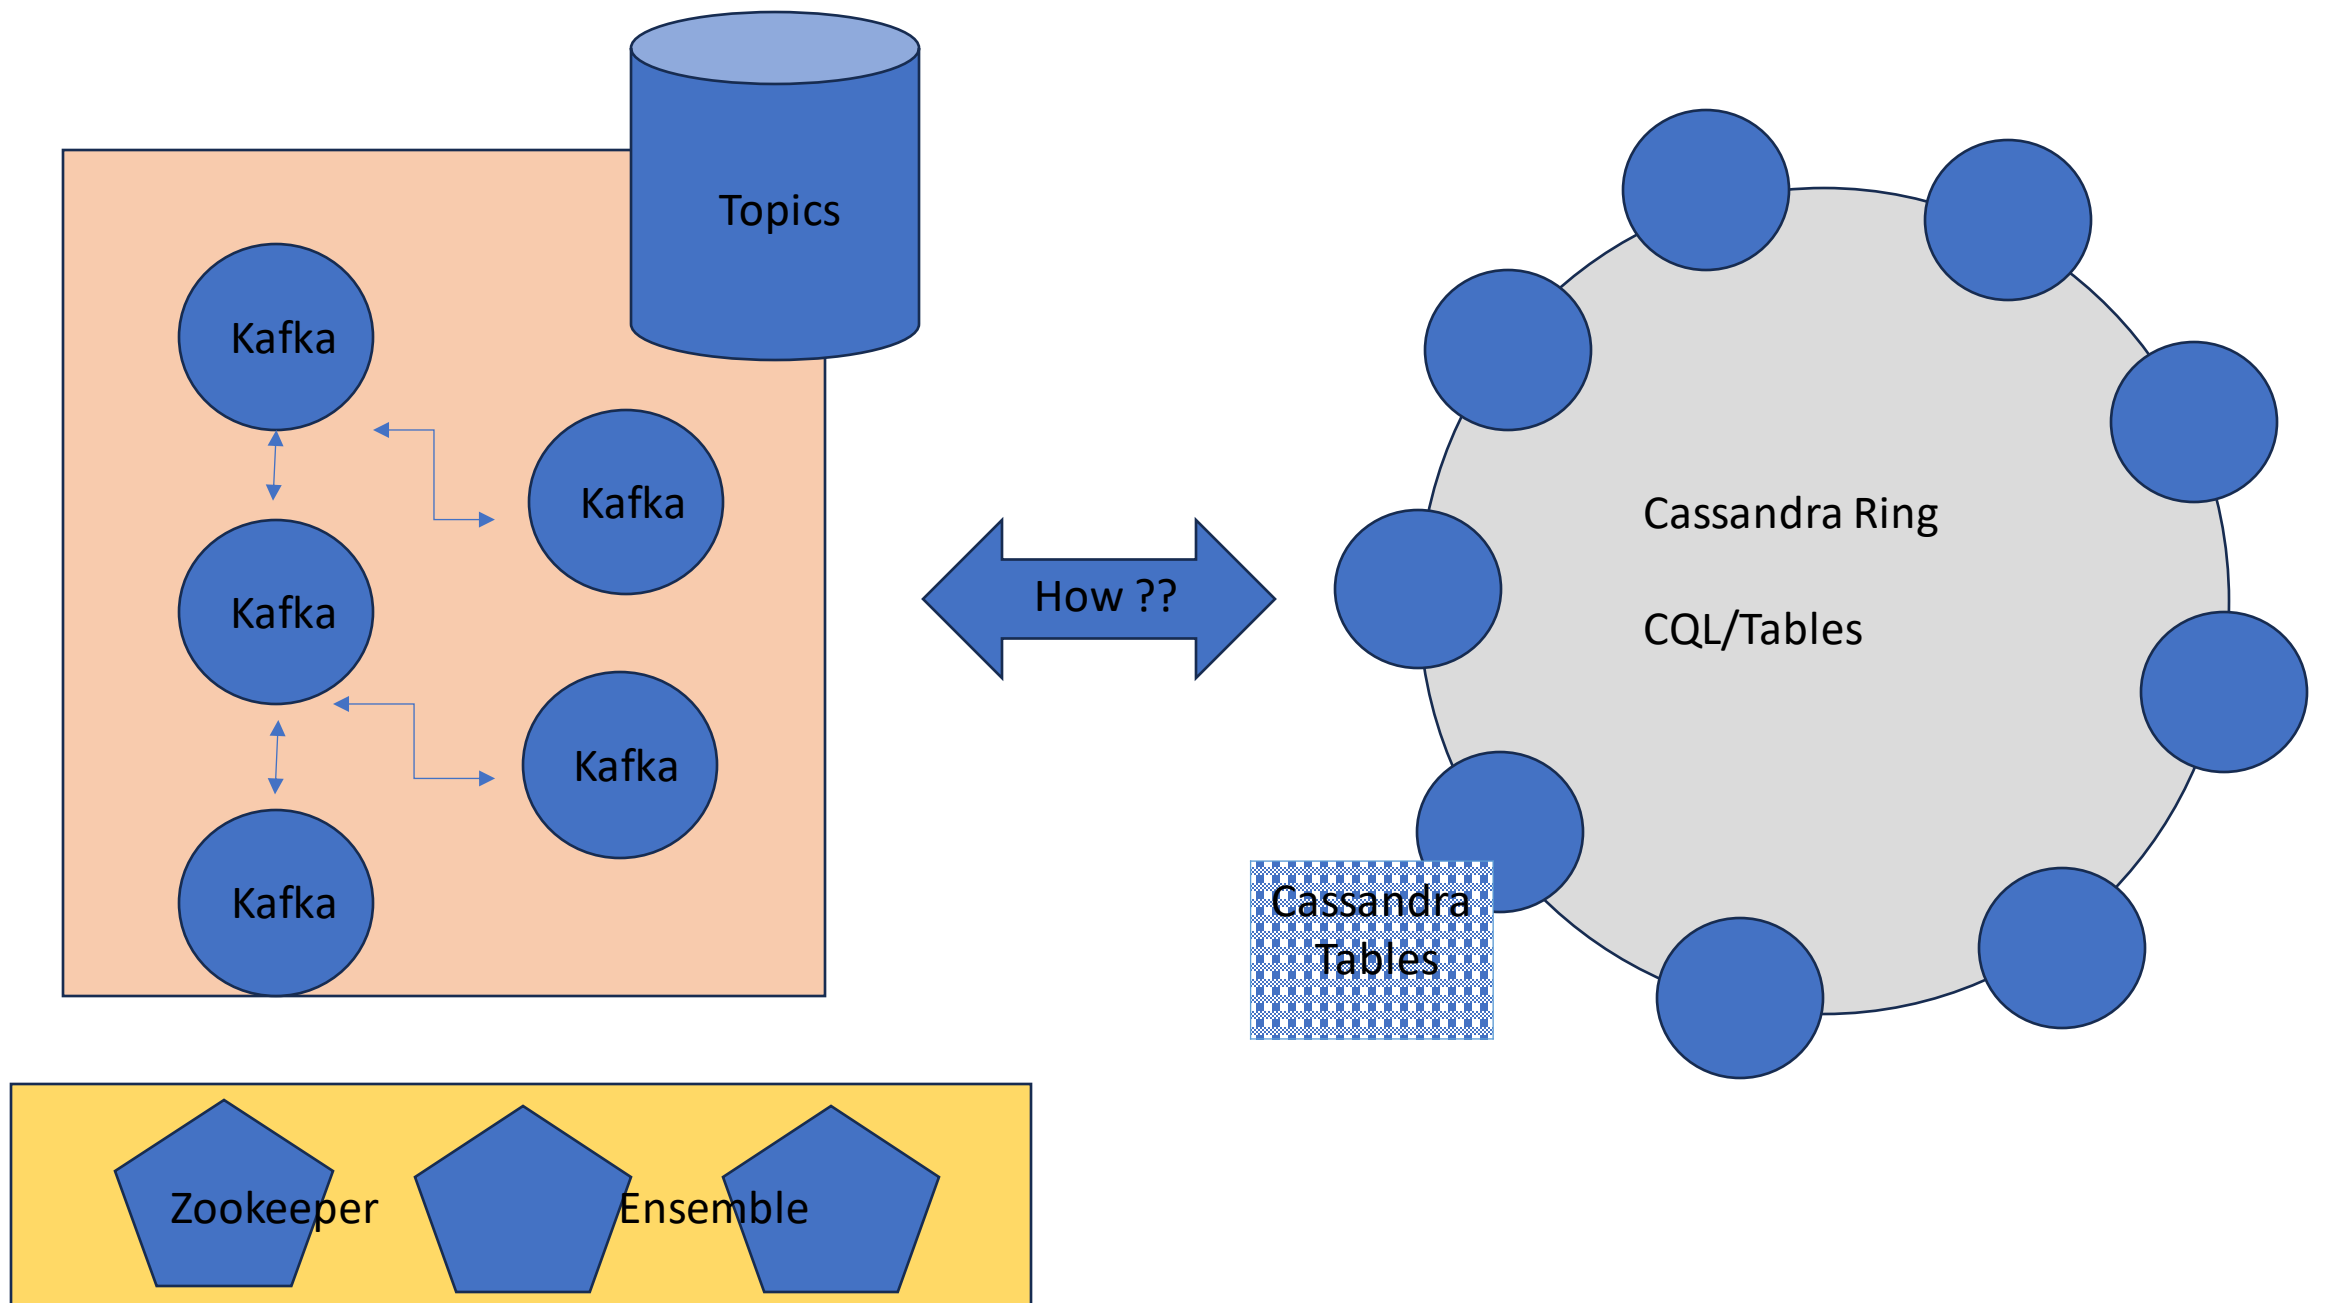
Two peas in a pod

Big Picture I

- Cassandra NoSQL performant database, still has one of the fastest I/O writes in the industry
- Kafka – zero copy technology, LinkedIn processes trillions of messages through their Kafka ecosystem
- Challenge how to get Kafka and Cassandra play nicely with each other
- **Complex scenario** – however, Kafka Connect to the rescue

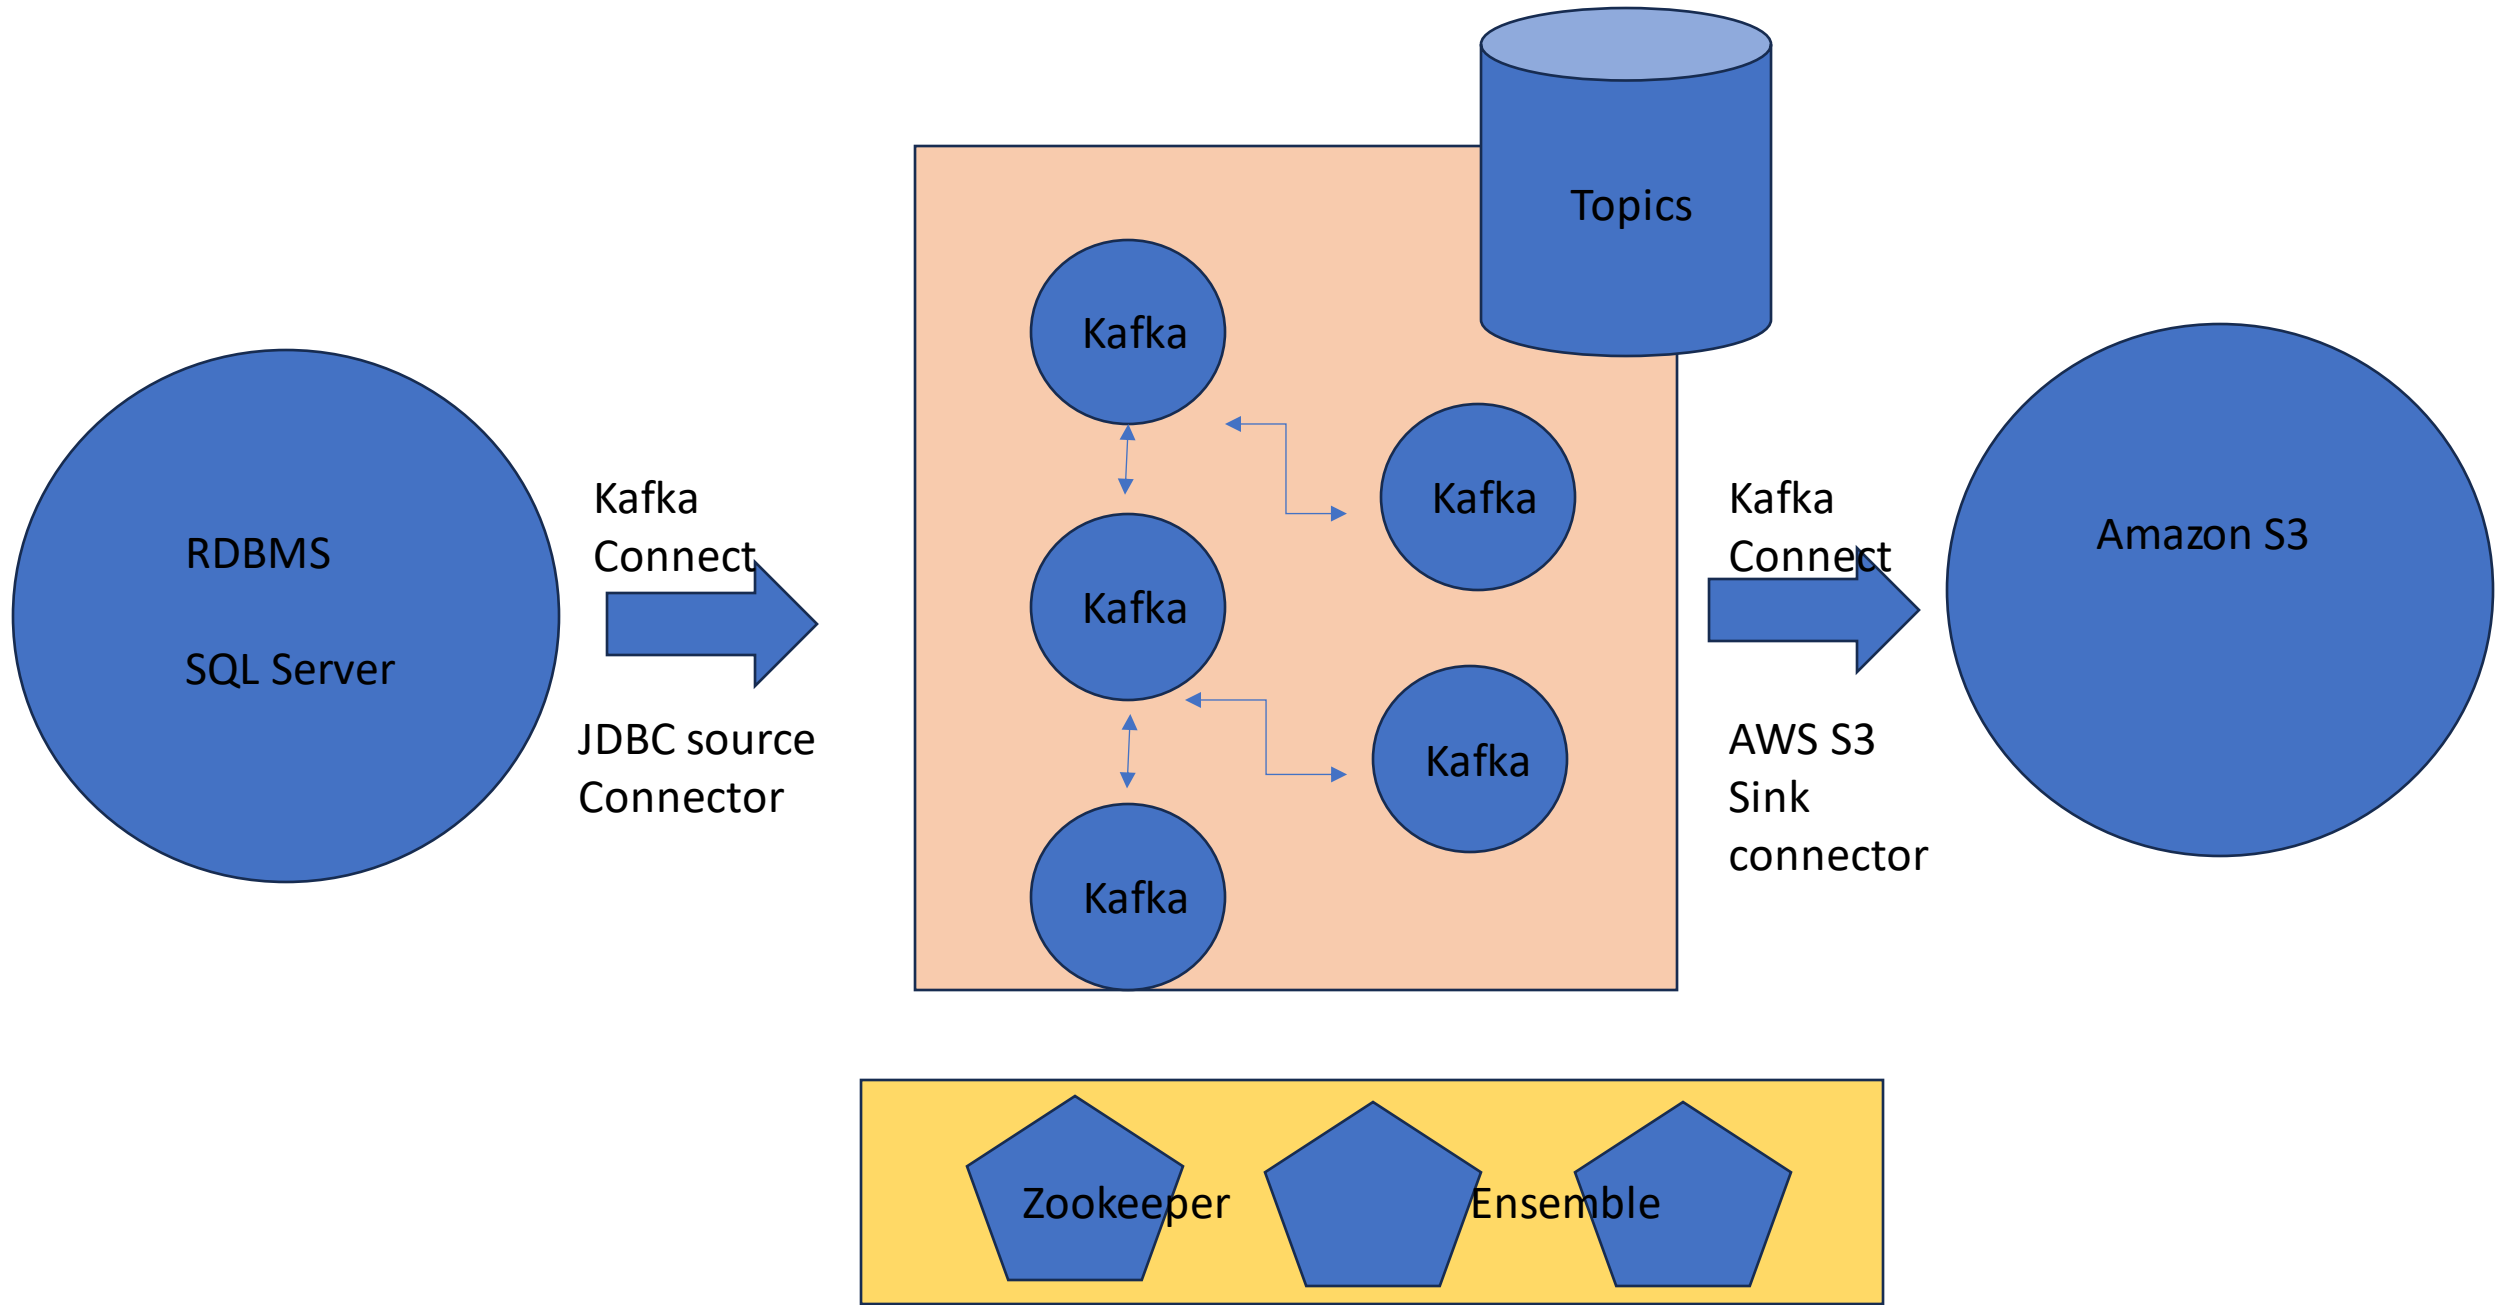
Big Picture II

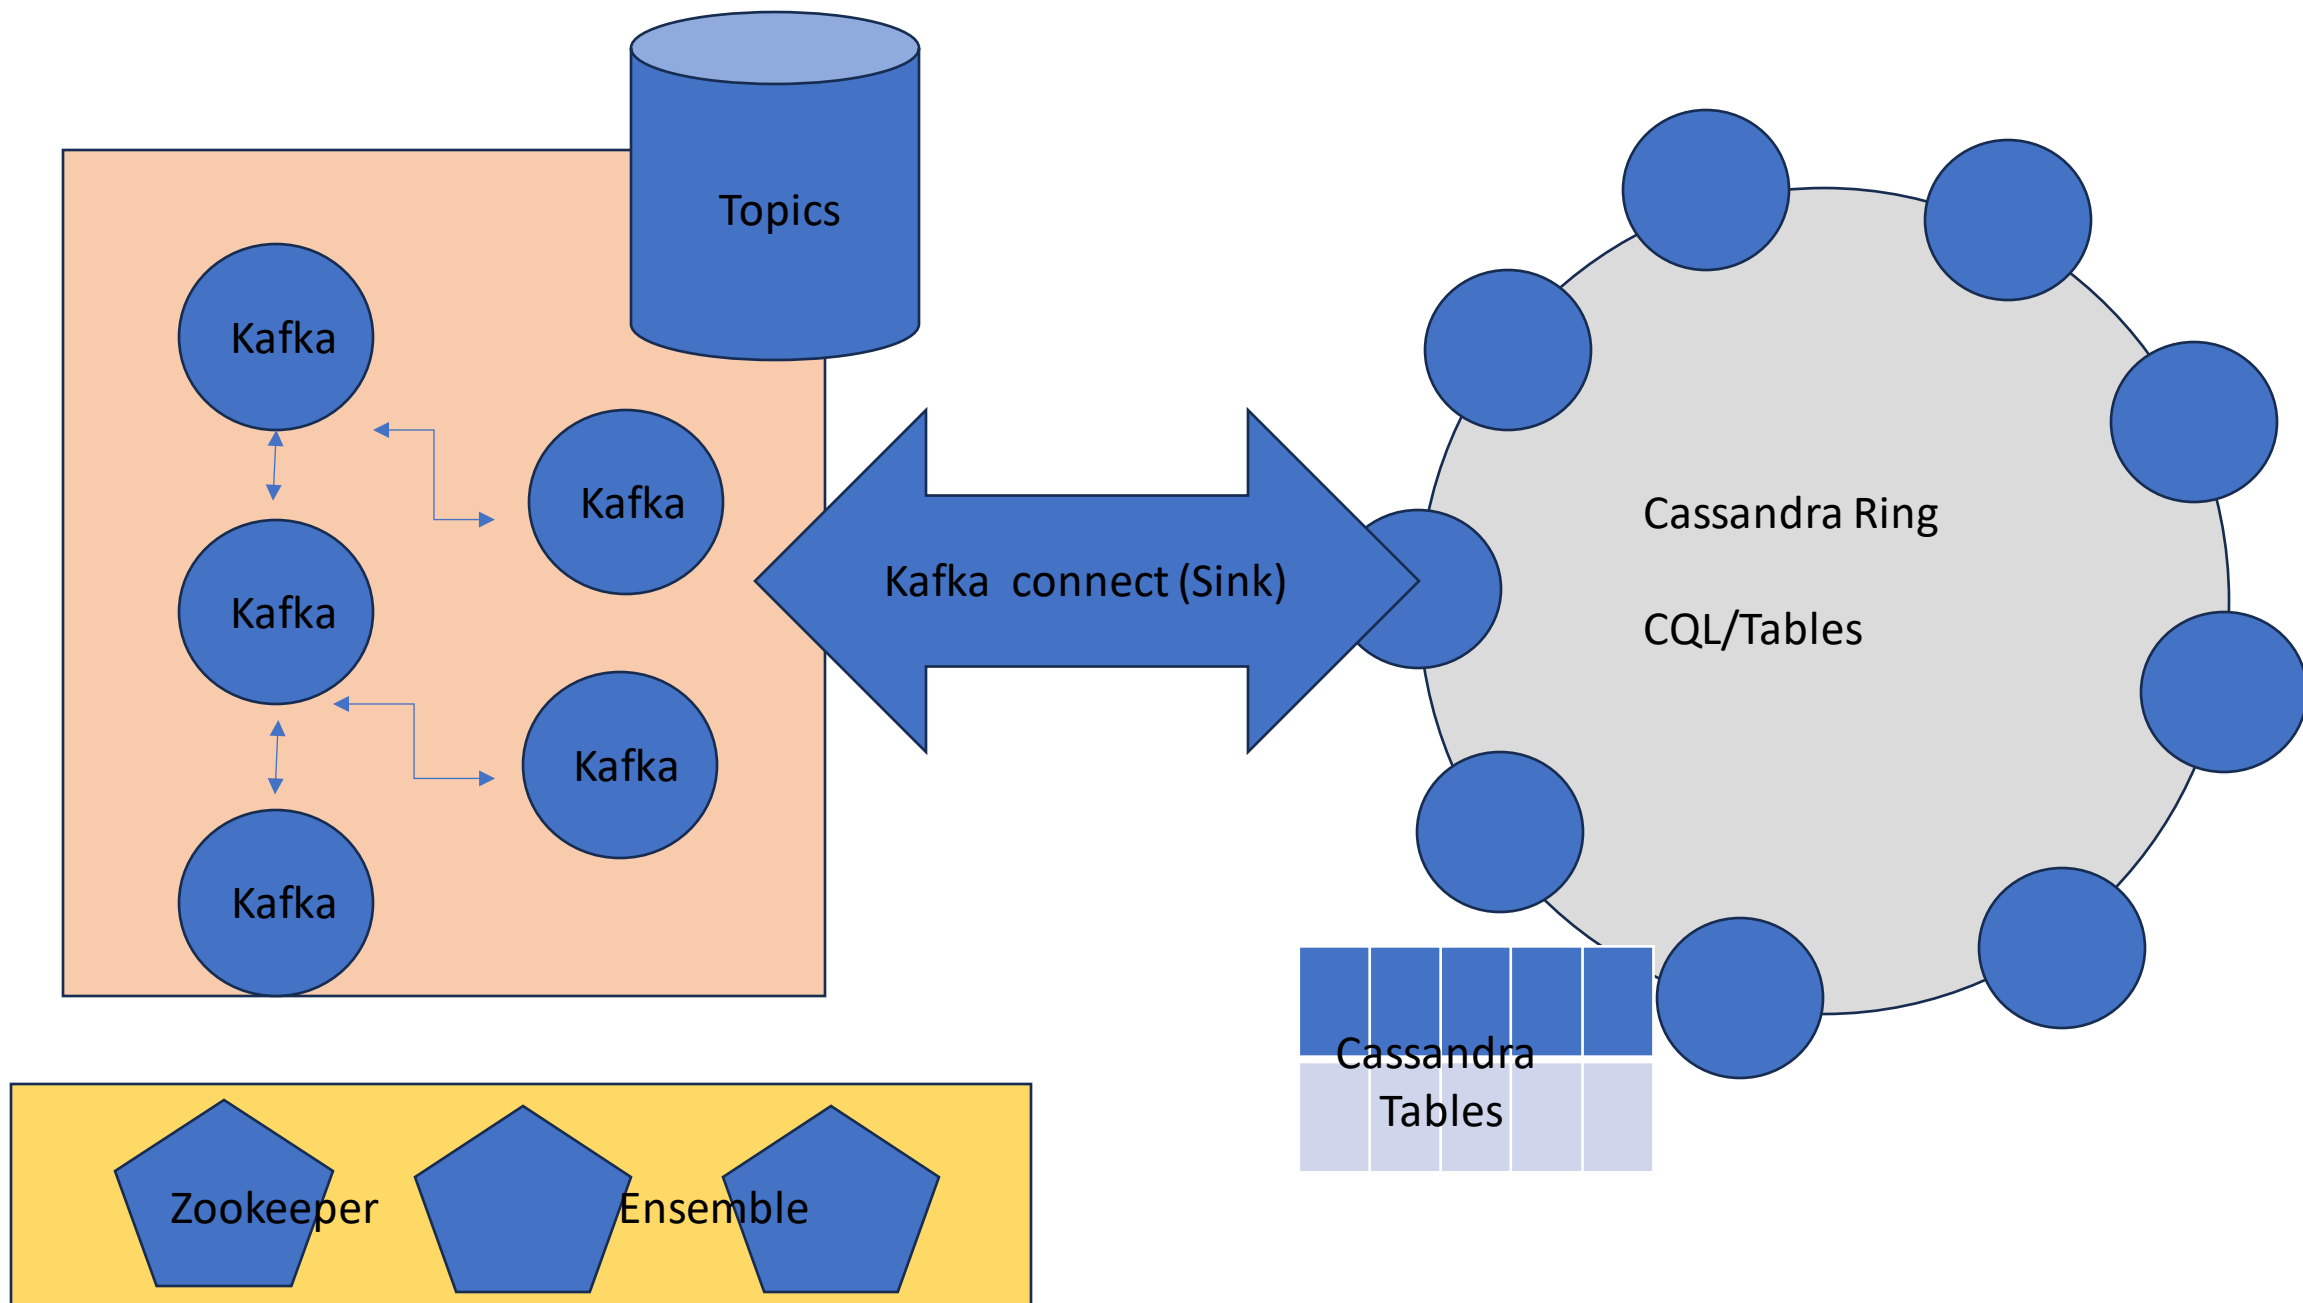
- Topics are the currency in Kafka
- Tables are the currency in Cassandra
- We want to do this
- Cassandra -> Kafka -> Cassandra
- Cassandra -> Kafka Connect -> Kafka -> Kafka Connect -> Cassandra
- Kafka Connect is the digital glue
- We will examine Kafka -> Kafka Connect (Sink) -> Cassandra
- Cassandra -> Kafka Connect(Source) -> Kafka (2nd half of slide deck)



Man, this is complex

Yes and No KAFKA CONNECT TO THE RESCUE






Kafka Connect(Sink)

- How much technology ? It is the properties file, stupid !!!
- Good news, no Java coding, just config files
- Two flavors – standalone and distributed
- Standalone good for dev, qa1 and qa2
- Distributed good for staging and prod
- Question – how do we go from Kafka topics to Cassandra tables
- All we have are 5 Kafka APIs
- Kafka Connect(Sink) is a Swiss Army knife – a smart toolset

Data Impedance Kafka vs Cassandra

I understand key value pairs, commit logs, segments and partitions




??????????

What to do



Cassandra Tables



Cassandra Tables



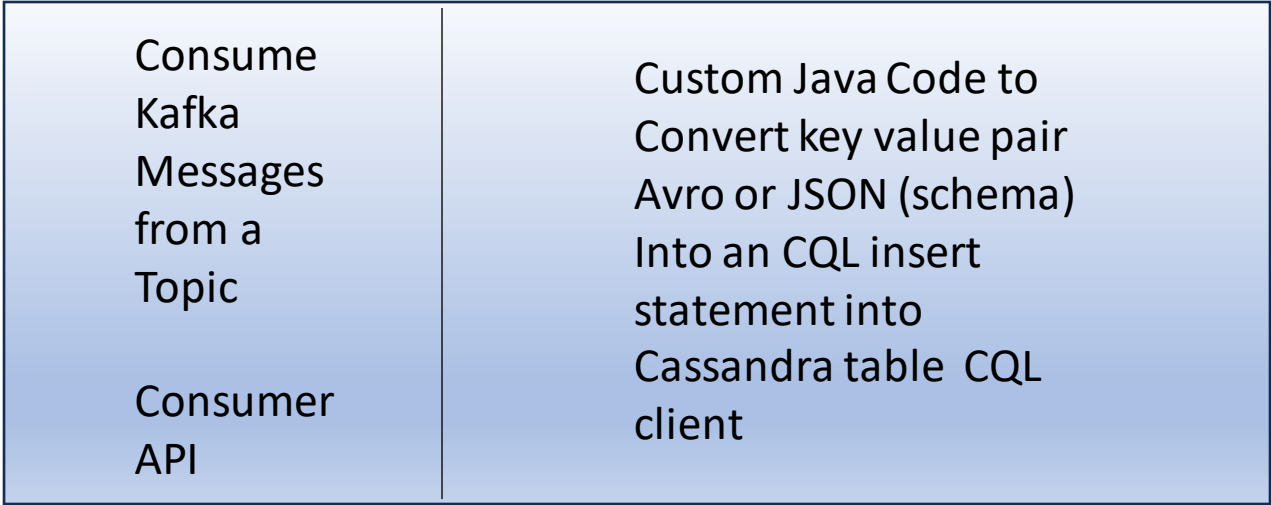


I understand nodes, CQL and CRUD statements

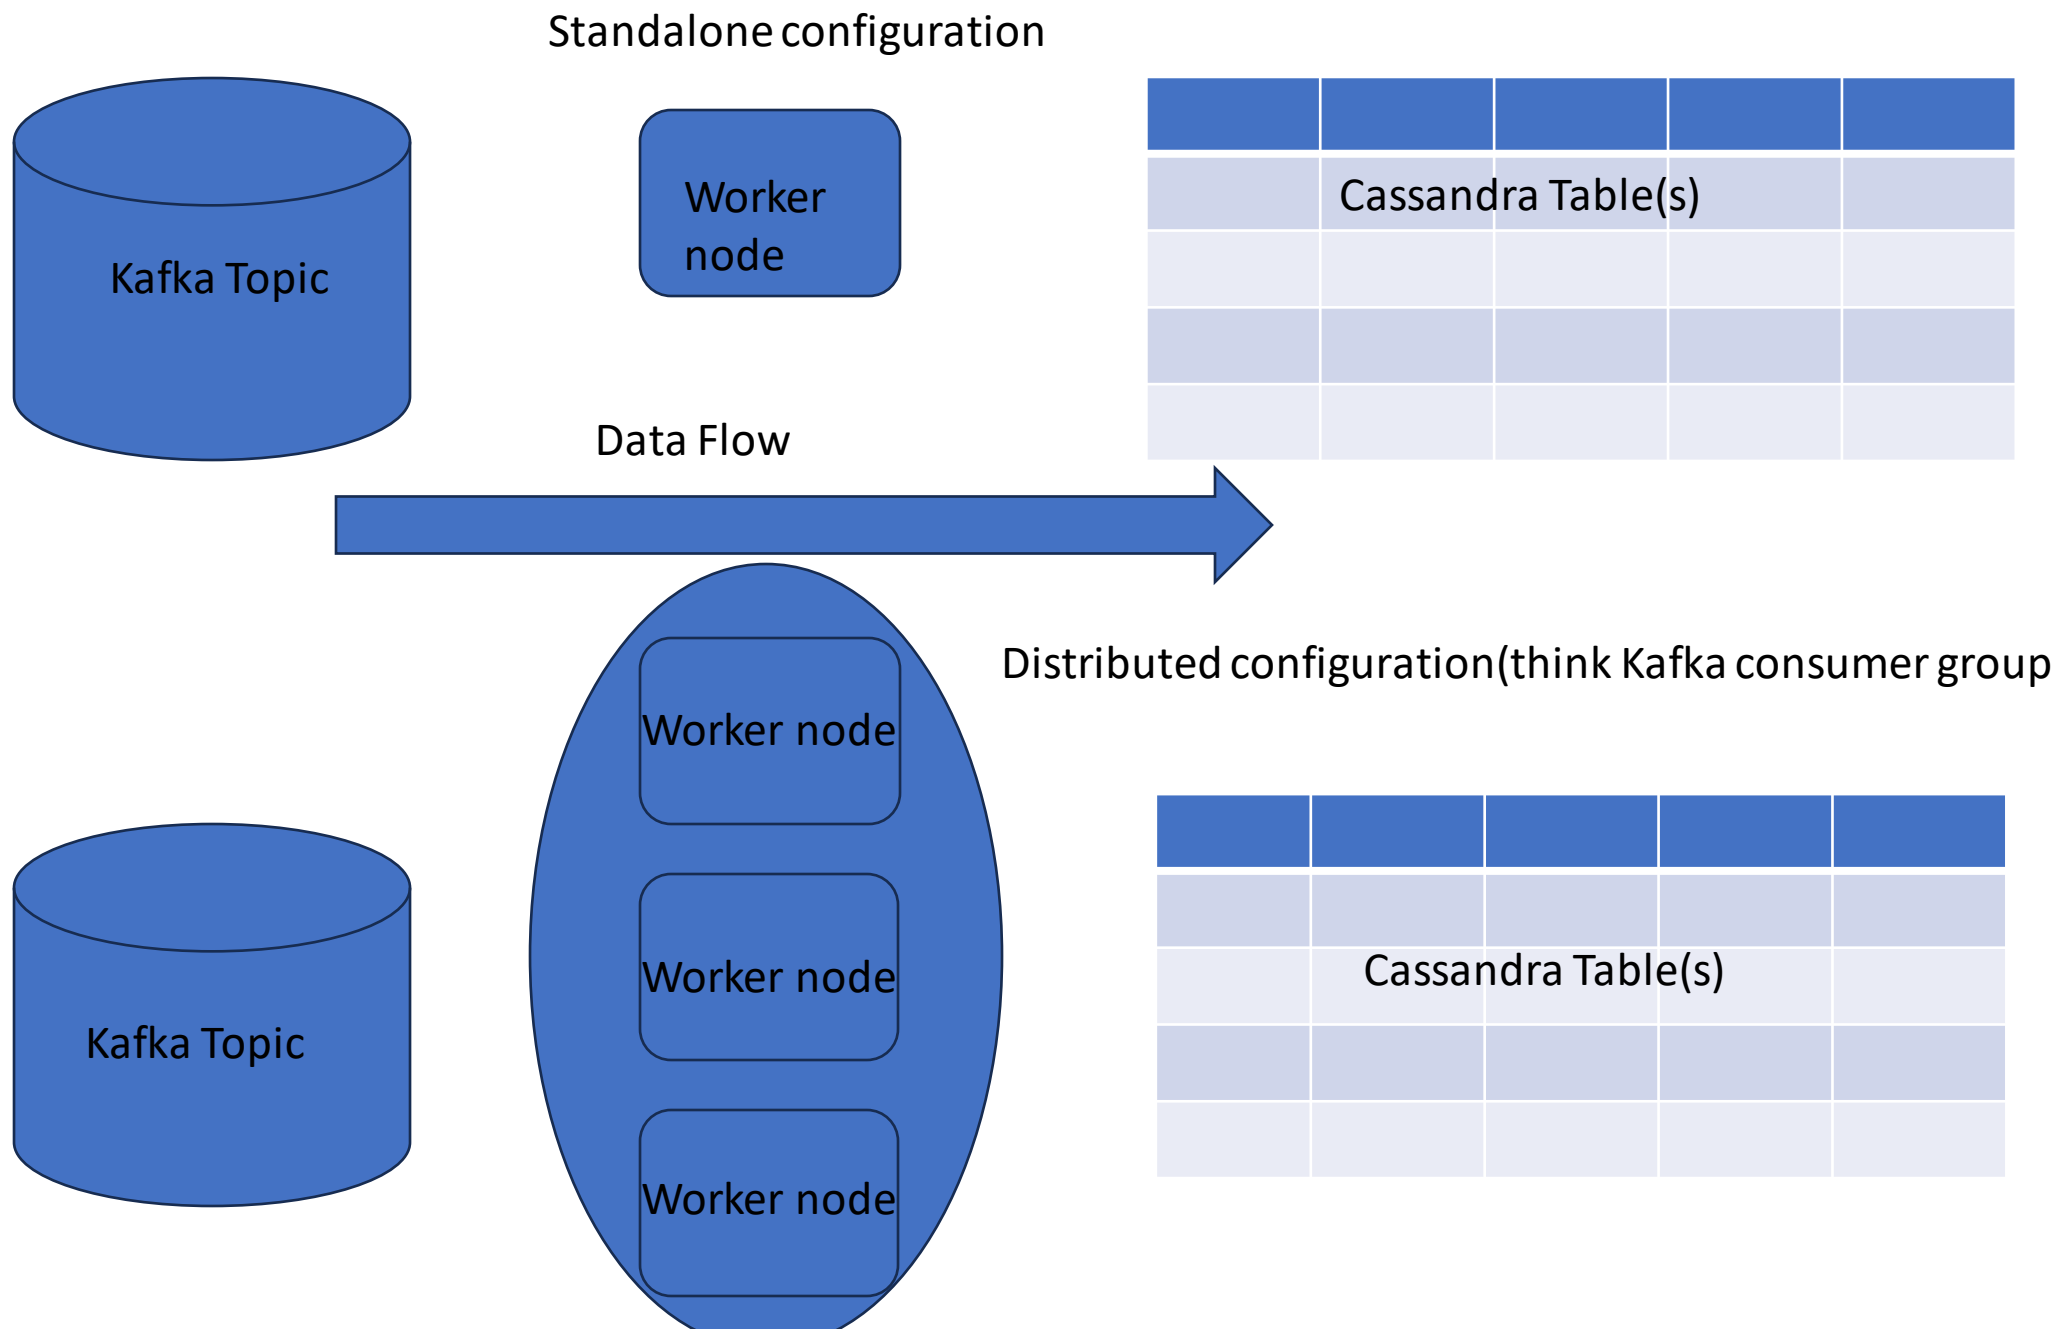
Kafka Connect
Worker
Node(standalone)

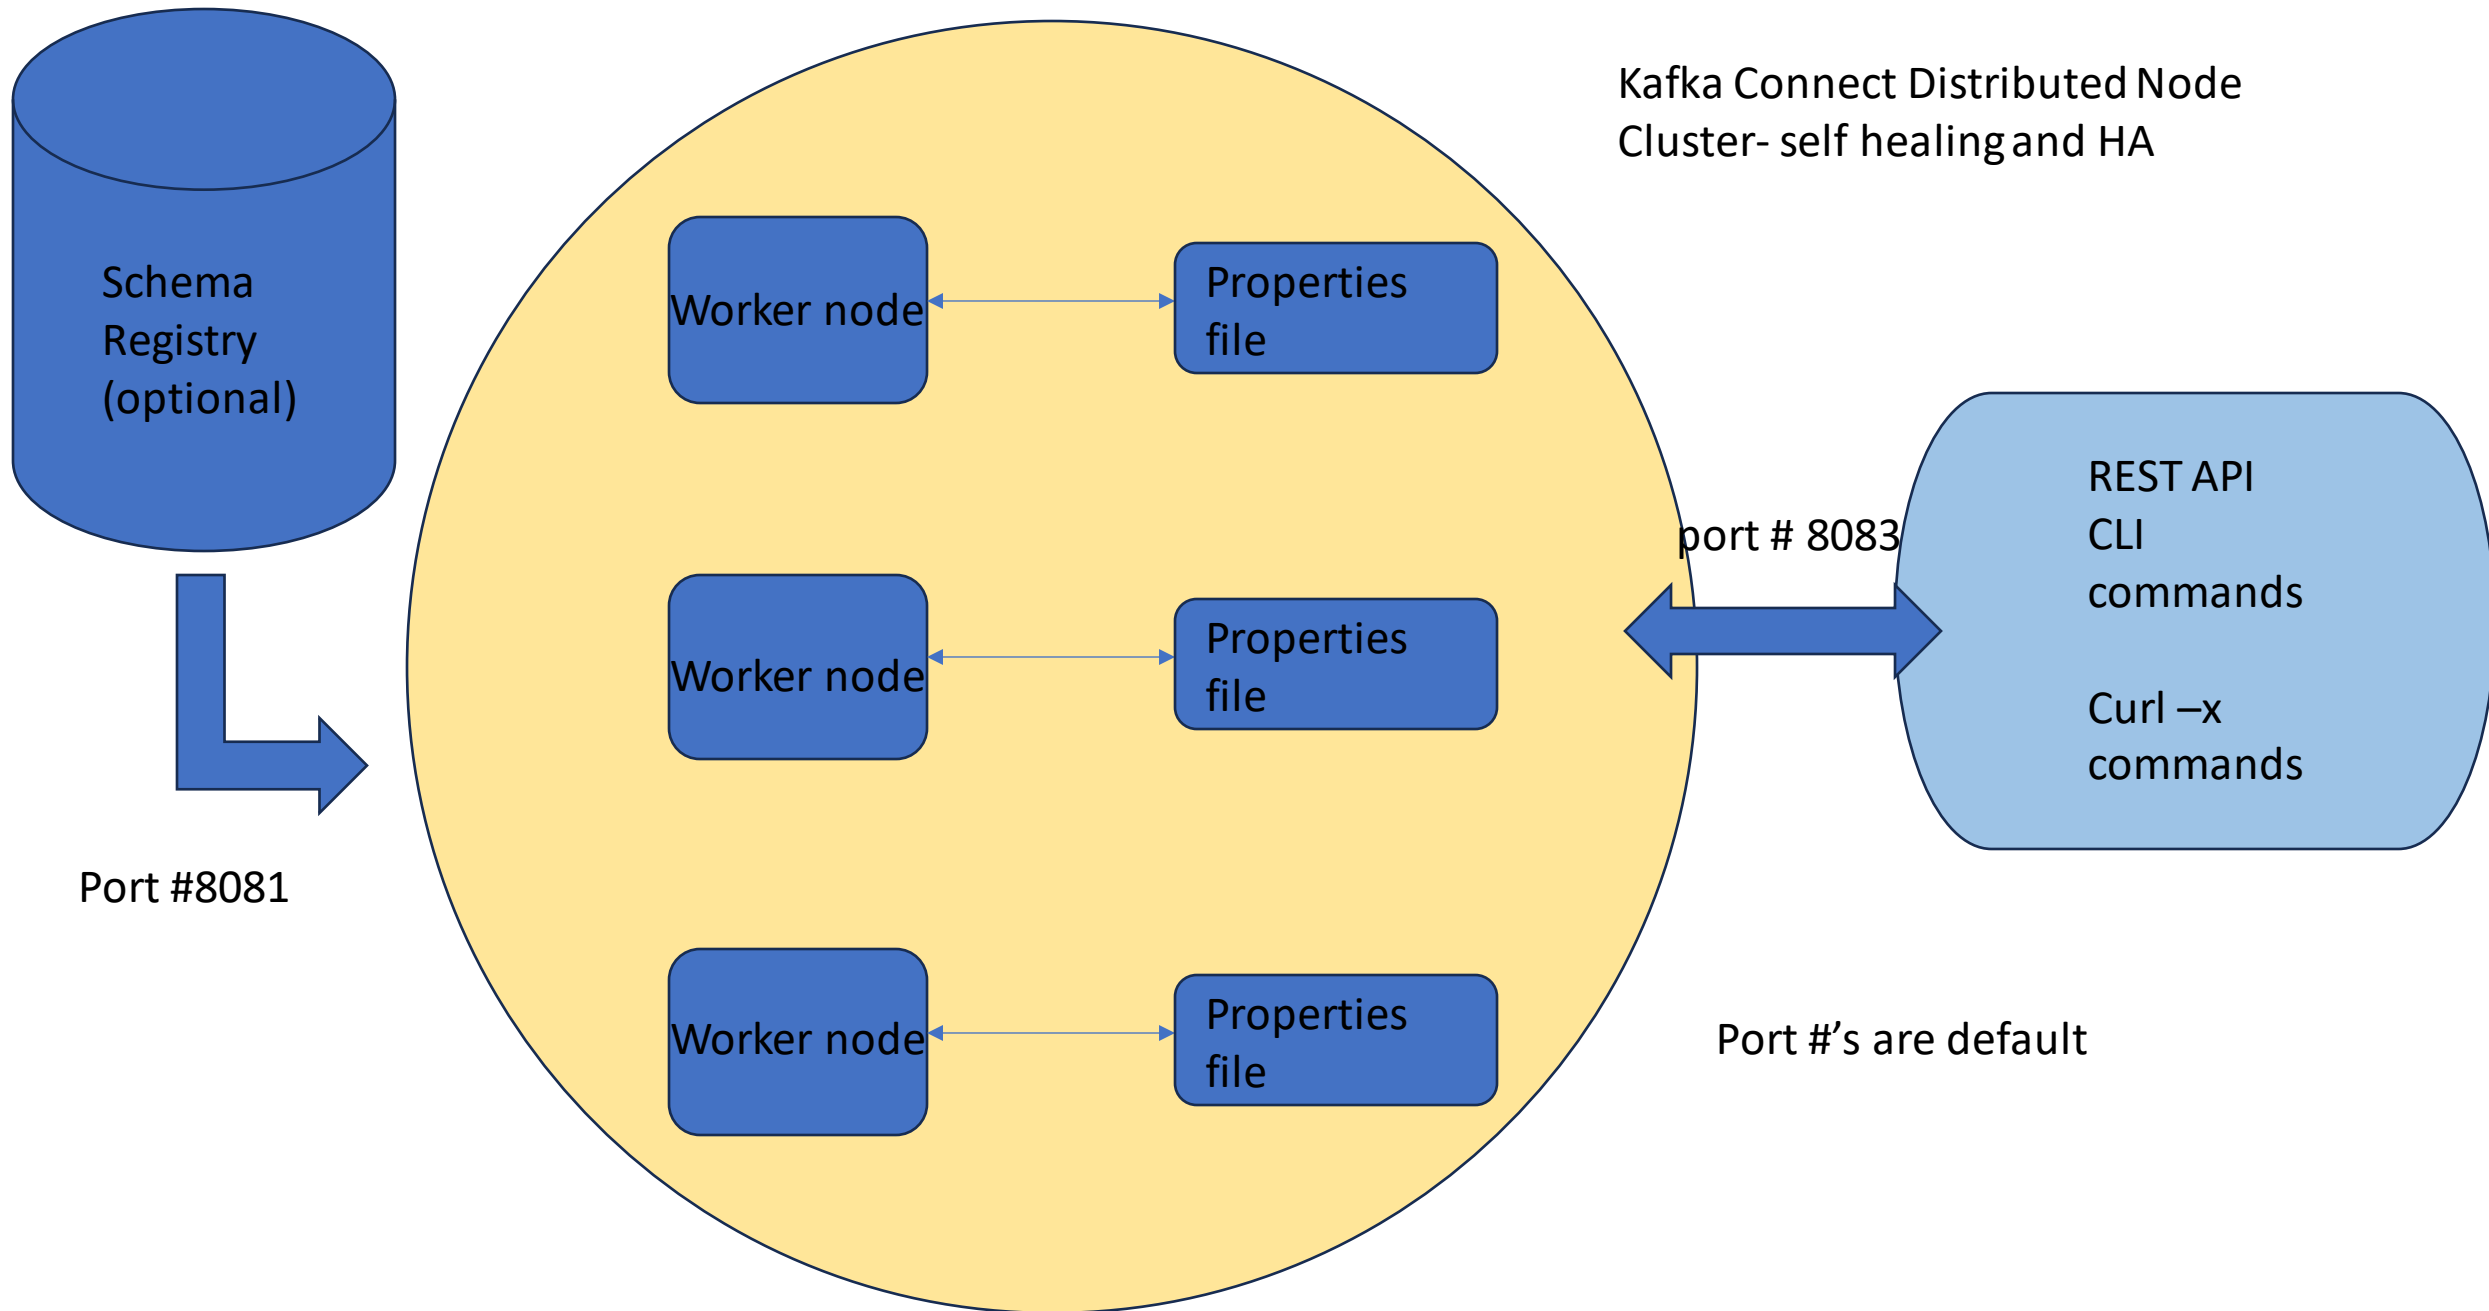
Kafka Connect Jar File :: plug-in

Direction of Data Flow (unidirectional)



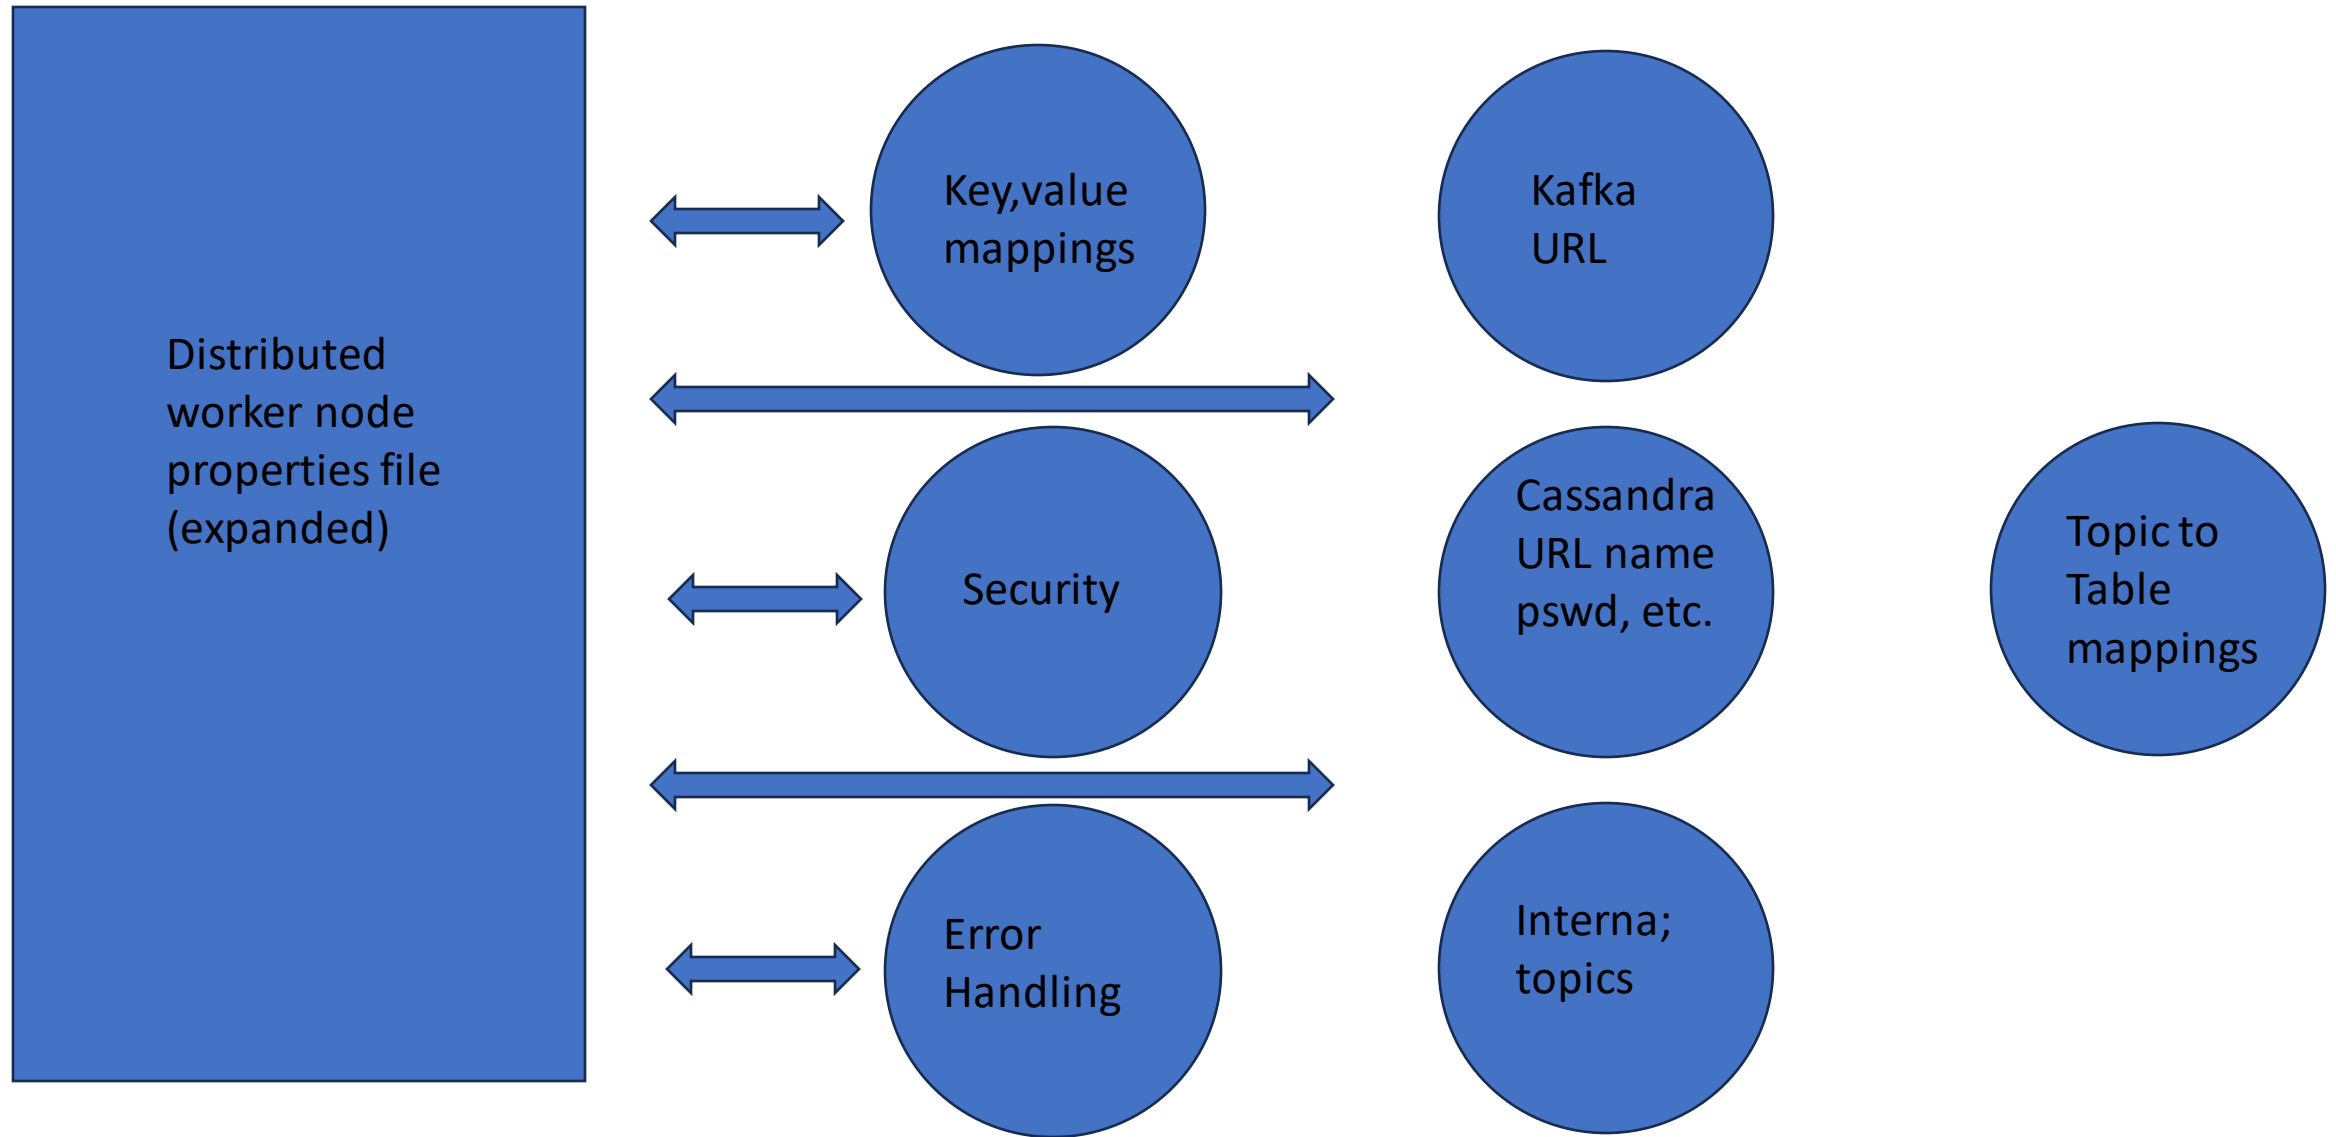
The Kafka Connect
Properties file runs
The show; details
to follow





Deep dive into worker node

- one property file per worker node
- principal use of distributed configuration – load balancing of data traffic ingestion into the Cassandra ring; don't worry Cassandra can keep up
- remember the secret sauce is the plug-in jar file does tge heavy lifting
- Property file points to Kafka broker list; Cassandra port/url string(aha!)
- Property file controls mapping of Kafka topics to Cassandra tables
- One worker node per Cassandra node for Cassandra keyspace
- Good news, no Java or Python scripting, some complexity in config file(s)



Where to get the jar files(plugin)

- go to **<http://www.confluent.io/hub/>**
- in the search bar, type Cassandra Sink Connector; hit enter key
- follow the directions for the download(s)
- save the jar file(plugin) copy the file into the worker node directory
- it is fully supported by Confluent
- don't be alarmed the zip file has a /lib directory of about 30+ jar files; transfer everything
- copy the entire directory (copy this library into the worker node directory, e.g., /usr/kafka/plugin)

# General worker configuration	Name of the worker node ,group id
# Kafka and Cassandra converters configuration	Set up the key value converters
# Set the plugin path to the directory containing Cassandra Sink Connector JAR	Point to the magic sauce plugin jar filesssserrrr
# Cassandra Sink Connector configuration for Worker 1	Self explanatory
# Cassandra connection details for Worker 1	Pointers to the key space, username,password, port, etc.
# Kafka topic to Cassandra table mapping for Worker 1	Mapping from Kafka topic(s) to Cassandra table(s)
#Security	SASL/SSL/truststores(tbd)

Name=worker-1

Bootstrap.servers=kafka-broker-1:9092, kafka-broker-2:9092;kafka-broker-3:9092

group-.id=connect-cluster

#Kafka and Cassandra converters configuration

key.converter=org.apache.kafka.connect.json.JsonConverter

value.converter=org.apache.kafka.connect.json.JsonConverter

Key.converter.schemas.enable=false

Value.converter.schemas.enable=false

#set the plugin path to the directory containing Cassandra Sink Connector

Plug-in.path=/path/to/connectors

#Cassandra Sink Connector configuration for worker-1

Connector.class=io.confluent.connect.Cassandra.CassandraSinkConnector

tasks=1

#Cassandra connection details for worker -1

Cassandra.connection.host=cassandra-node-1

Cassandra.connection.port=9042

Cassandra.connection.username=my-username

Cassandra.connection.password=my-password

Cassandra.keyspace=my_keyspace

#Kafka topic to Cassandra table mapping for worker -1

topic.myworker_1-topic_1=Cassandra_table_1

topic.myworker_1-topic_2=Cassandra_table_2

topic.myworker_1-topic_3=Cassandra_table_3

Name=worker-2

Bootstrap.servers=kafka-broker-1:9092, kafka-broker-2:9092;kafka-broker-3:9092

group-.id=connect-cluster

#Kafka and Cassandra converters configuration

key.converter=org.apache.kafka.connect.json.JsonConverter

value.converter=org.apache.kafka.connect.json.JsonConverter

Key.converter.schemas.enable=false

Value.converter.schemas.enable=false

#set the plugin path to the directory containing Cassandra Sink Connector

Plug-in.path=/path/to/connectors

#Cassandra Sink Connector configuration for worker-2

Connector.class=io.confluent.connect.Cassandra.CassandraSinkConnector

tasks=1

#Cassandra connection details for worker -1

Cassandra.connection.host=cassandra-node-1

Cassandra.connection.port=9042

Cassandra.connection.username=my-username

Cassandra.connection.password=my-password

Cassandra.keyspace=my_keyspace

#Kafka topic to Cassandra table mapping for worker -1

topic.myworker_2-topic_1=Cassandra_table_1

topic.myworker_2-topic_2=Cassandra_table_2

topic.myworker_2-topic_3=Cassandra_table_3

Name=worker-3
Bootstrap.servers=kafka-broker-1:9092, kafka-broker-2:9092;kafka-broker-3:9092
group-.id=connect-cluster
#Kafka and Cassandra converters configuration
key.converter=org.apache.kafka.connect.json.JsonConverter
value.converter=org.apache.kafka.connect.json.JsonConverter
Key.converter.schemas.enable=false
Value.converter.schemas.enable=false

#set the plugin path to the directory containing Cassandra Sink Connector
Plug-in.path=/path/to/connectors

#Cassandra Sink Connector configuration for worker-3
Connector.class=io.confluent.connect.Cassandra.CassandraSinkConnector
tasks=1

#Cassandra connection details for worker -1
Cassandra.connection.host=cassandra-node-1
Cassandra.connection.port=9042
Cassandra.connection.username=my-username
Cassandra.connection.password=my-password
Cassandra.keyspace=my_keyspace

#Kafka topic to Cassandra table mapping for worker -3

topic.myworker_3-topic_1=Cassandra_table_4

topic.myworker_3-topic_2=Cassandra_table_5

topic.myworker_3-topic_3=Cassandra_table_6

Cassandra Settings(Common)

#Cassandra connection details

cassandra.contact.points=localhost #IP address

cassandra.port=9042

cassandra.username=your username

cassandra.password=your password

cassandra.keyspace=your keyspace

cassandra.retryPolicy=DefaultRetryPolicy

cassandra.consistencyLevel=QUORUM

Error Handling(common)

#Error Handling

Errors.tolerance=all

Errors.log.enable=true

Errors.log.include.messages=true

Errors.deadletterqueue.topic.name=dlq-topic

Errors.deadletterqueue.contet.headers.enable=true

Errors.deadletterqueue.topic.replicationfactor=3

Errors.retry.delay.mas.ms=60000

Errors.retry.timeout=0

Internal topics Provisioning(Common)

#internal topics used by Kafka Connect

config.storage.topic=connect-configs

status.storage.topic=connect-status

offset.storage.topic=connect-offsets

offset.storage.partitions=1

offset.storage.replication.factor=1

#Note: all three properties files must share the same internal topic

names – they must be the same name(s)

Security I (common – tbd)

#enable SSL/TLS

Security.protocol=SSL

#location of the truststore containing trusted CA certificates

Ssl.truststore.location=/path/to/truststore.jks

Ssl.truststore.password=truststore password

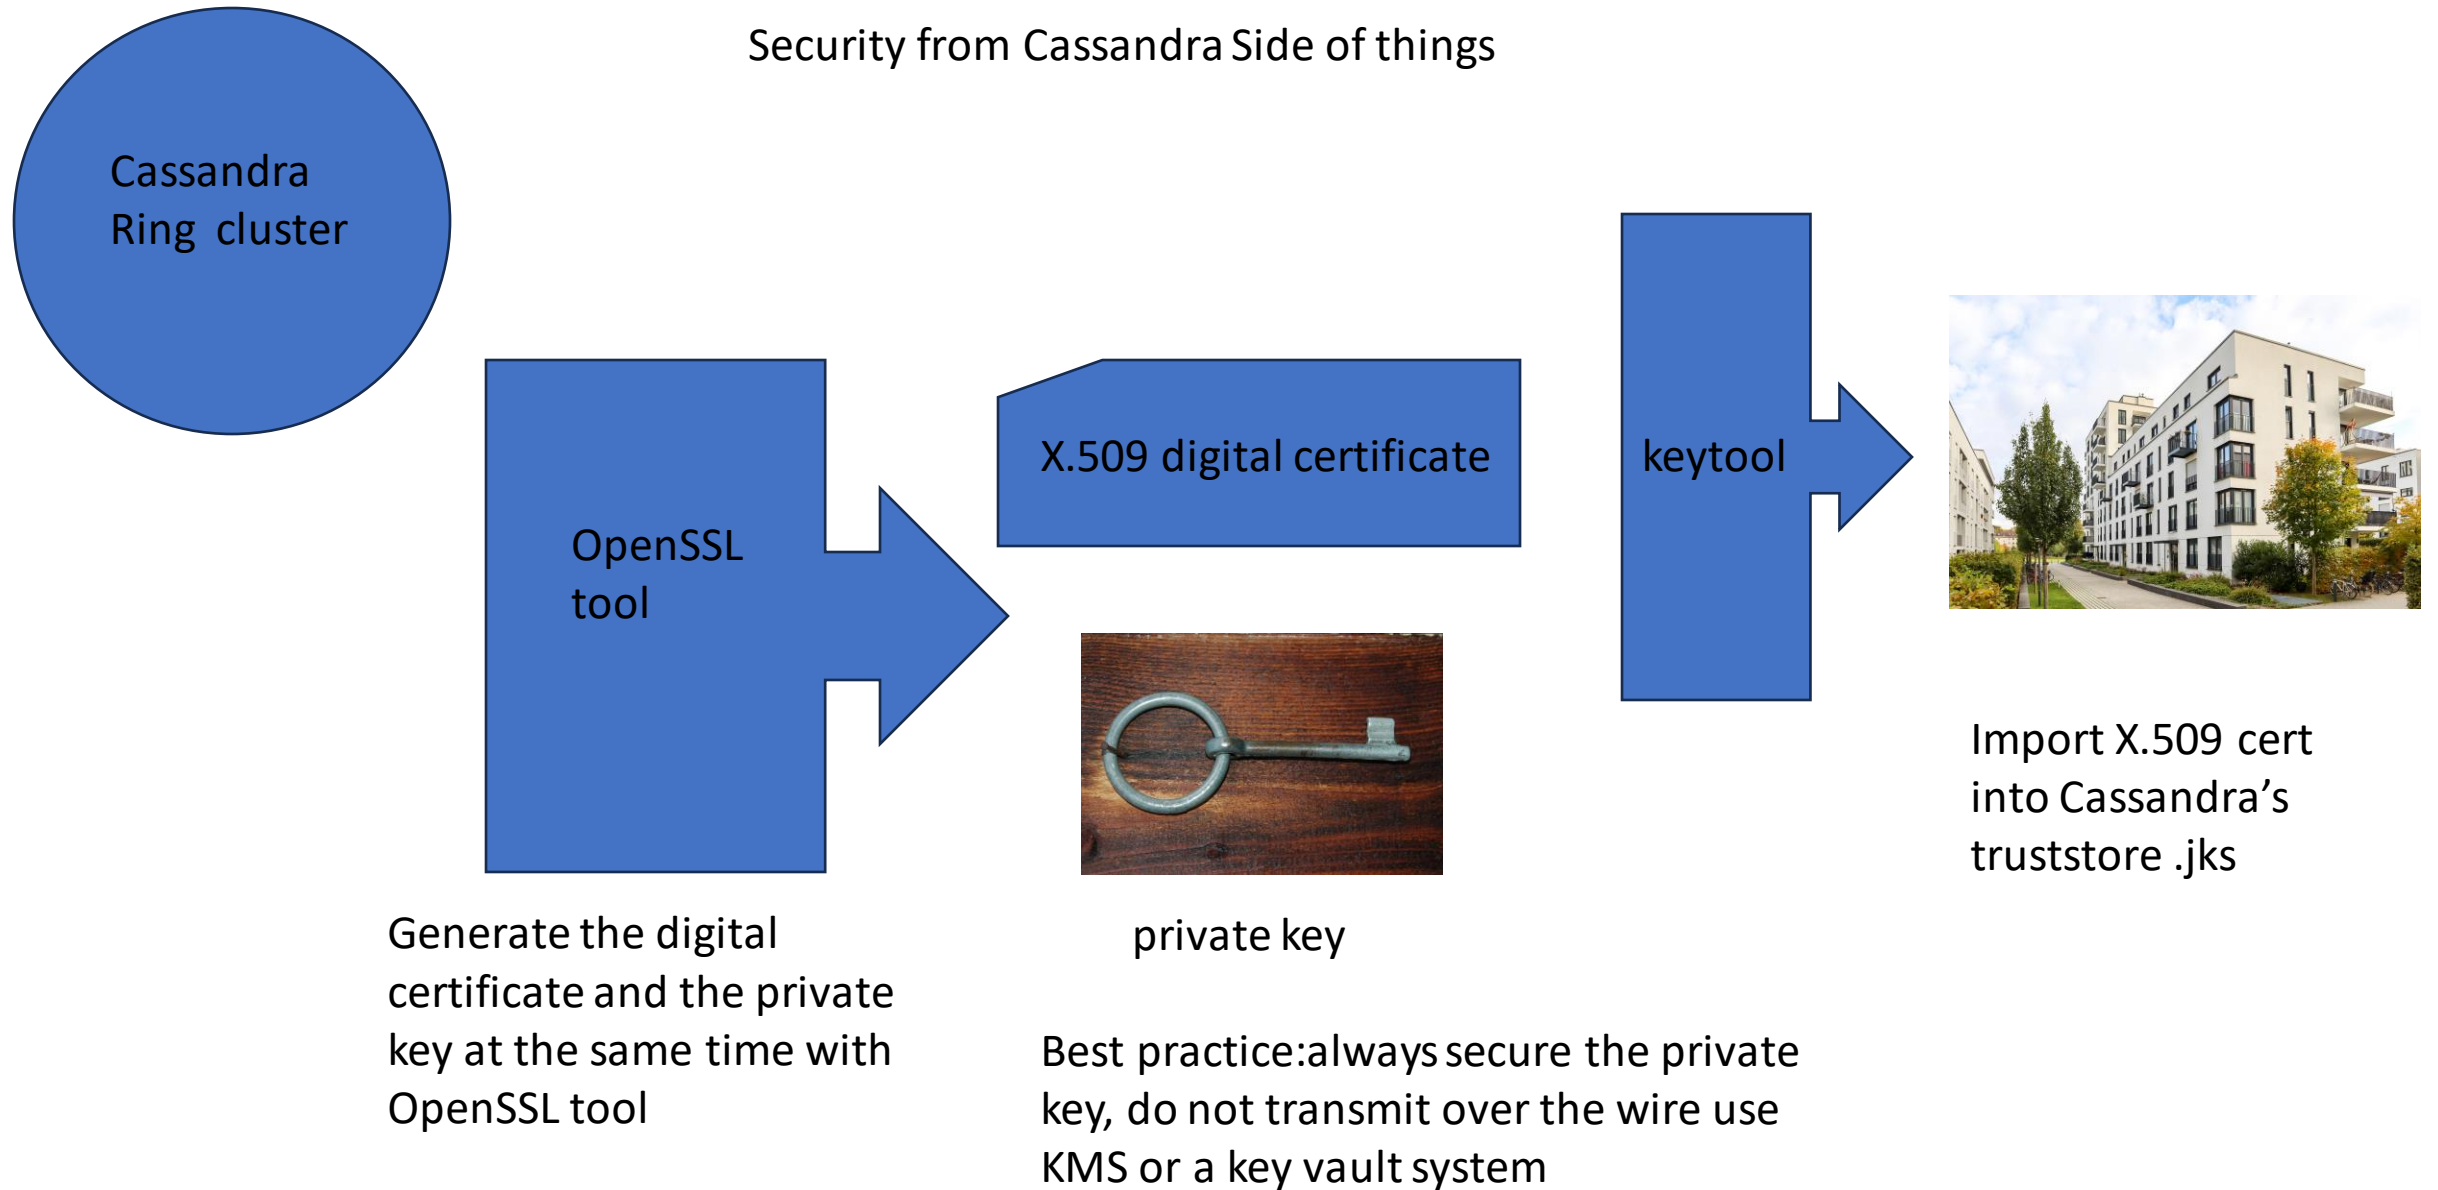
#if required



Ssl.keystore.location=/path/to/keystore.jks

Ssl.keystore.password=keystore password



Ssl.key.password= key password

Security from Cassandra Side of things








Generate a new keystore





Generate a new CSR

Certificate signing request



Sign the CSR
with CA
private key





Store signed CSR in Cassandra sink connector keystore
Strength of SSL/TLS: private key is never transmitted over the wire, it is secured
In the Kafka connect signed CSR with Cassandra's private key; On the Cassandra side of things it has the private key to decipher the signed CSR to match up and authenticate the sending party(Kafka connect side)

Security II (TBD) common

#enable SASL authentication

Security.protocol=SASL_PLAINTEXT

#SASL mechanism PLAIN, SCRAM-SHA-256,SCRAM-SHA-512

Sasl.mechanism=PLAIN

Sasl.jaas.config=org.apache.kafka.common.Security.plain.PlainLoginModule

#username and password

username=your username

Password=your password

p

REST API(Common)

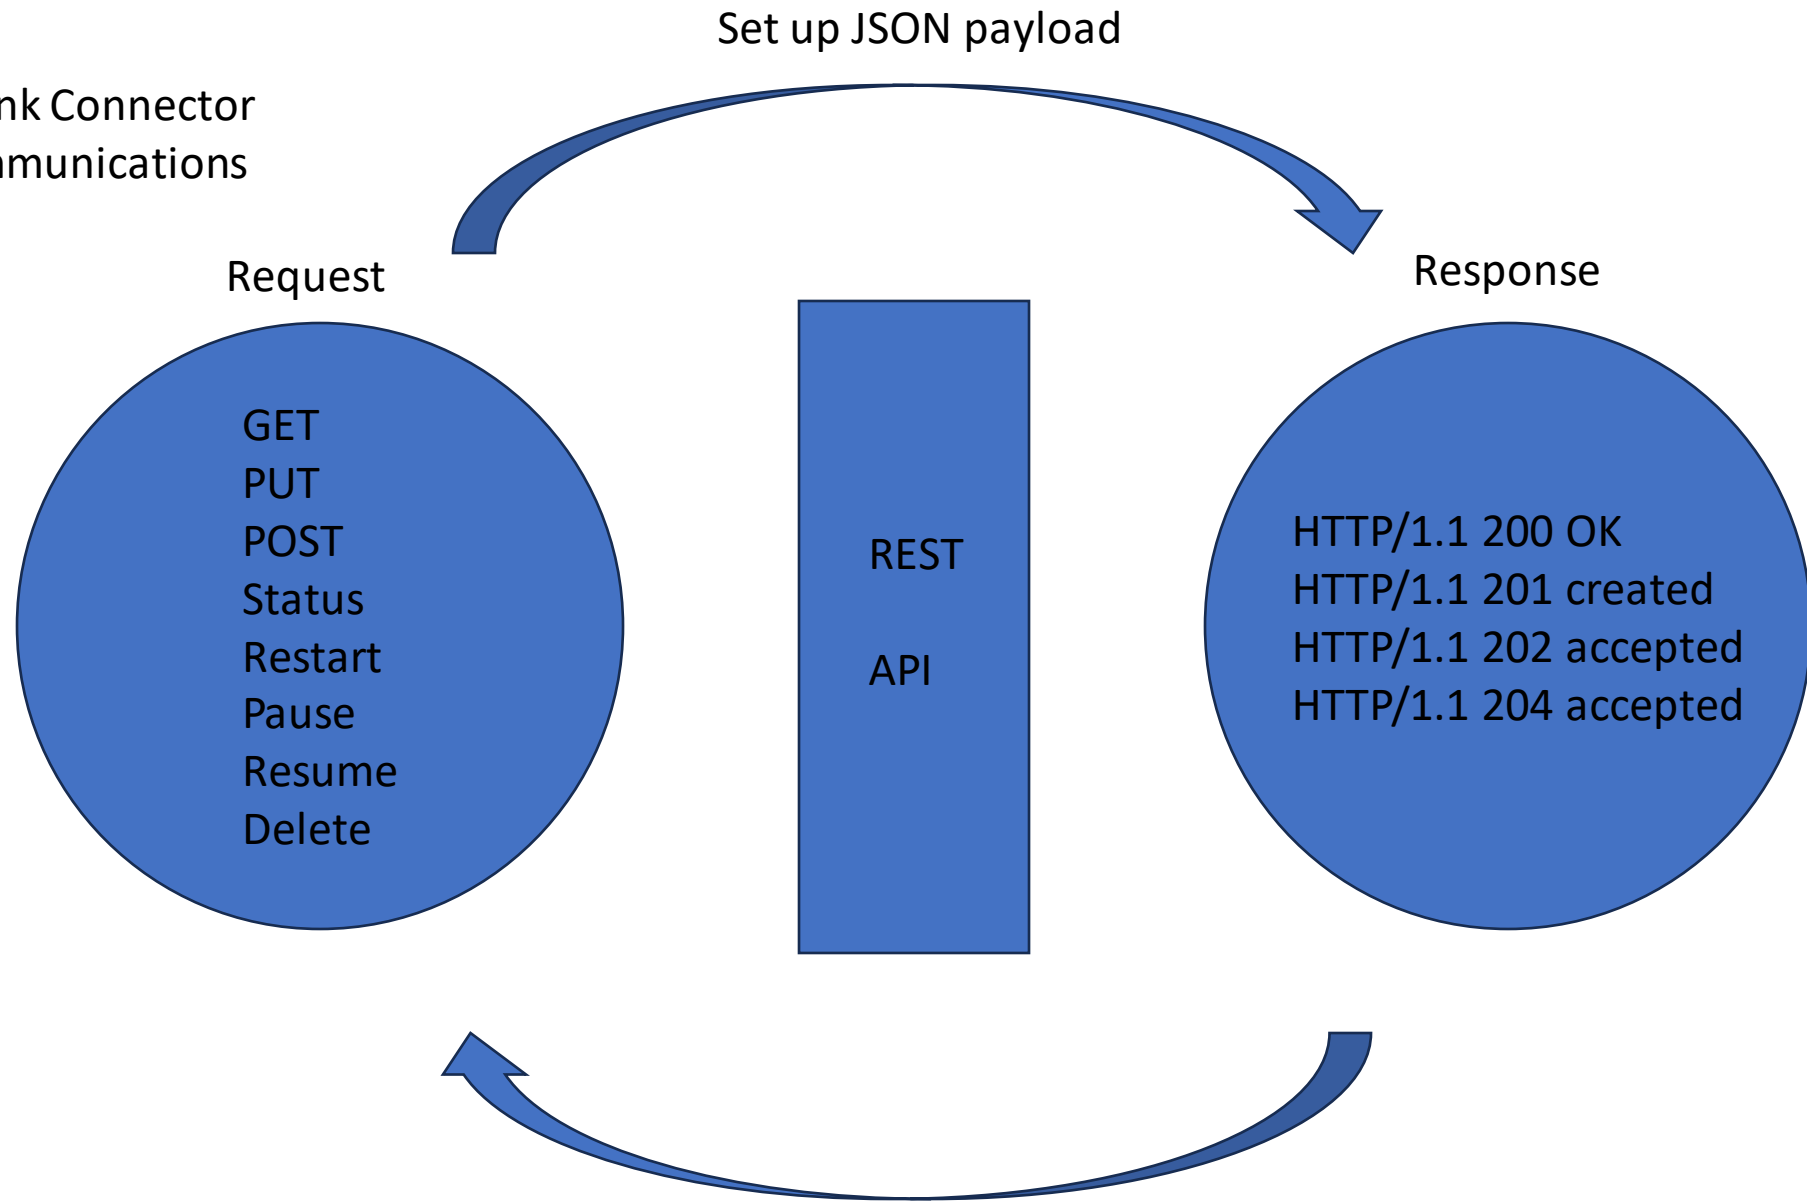
#rest api for distributed worker node

Rest.port=8083

Rest.advertised.host.name=localhost #real IP address

Rest.advertised.port=8083

Cassandra Sink Connector
Bilateral communications




We set up the distributed node properties

- time to execute !!!
- how do we run Kafka Connect Cassandra sink connector(s) ???
- two ways get the connect-distributd.sh shell script from Apache directly or standalone connect-standalone.sh script
- get the shell scripts from Confluent distribution subscription /bin
- **./bin/connect-distributed.sh config/connect-distributed.properties**
- connection-distribution properties is the “quarterback of this Kafka Connect ecosystem

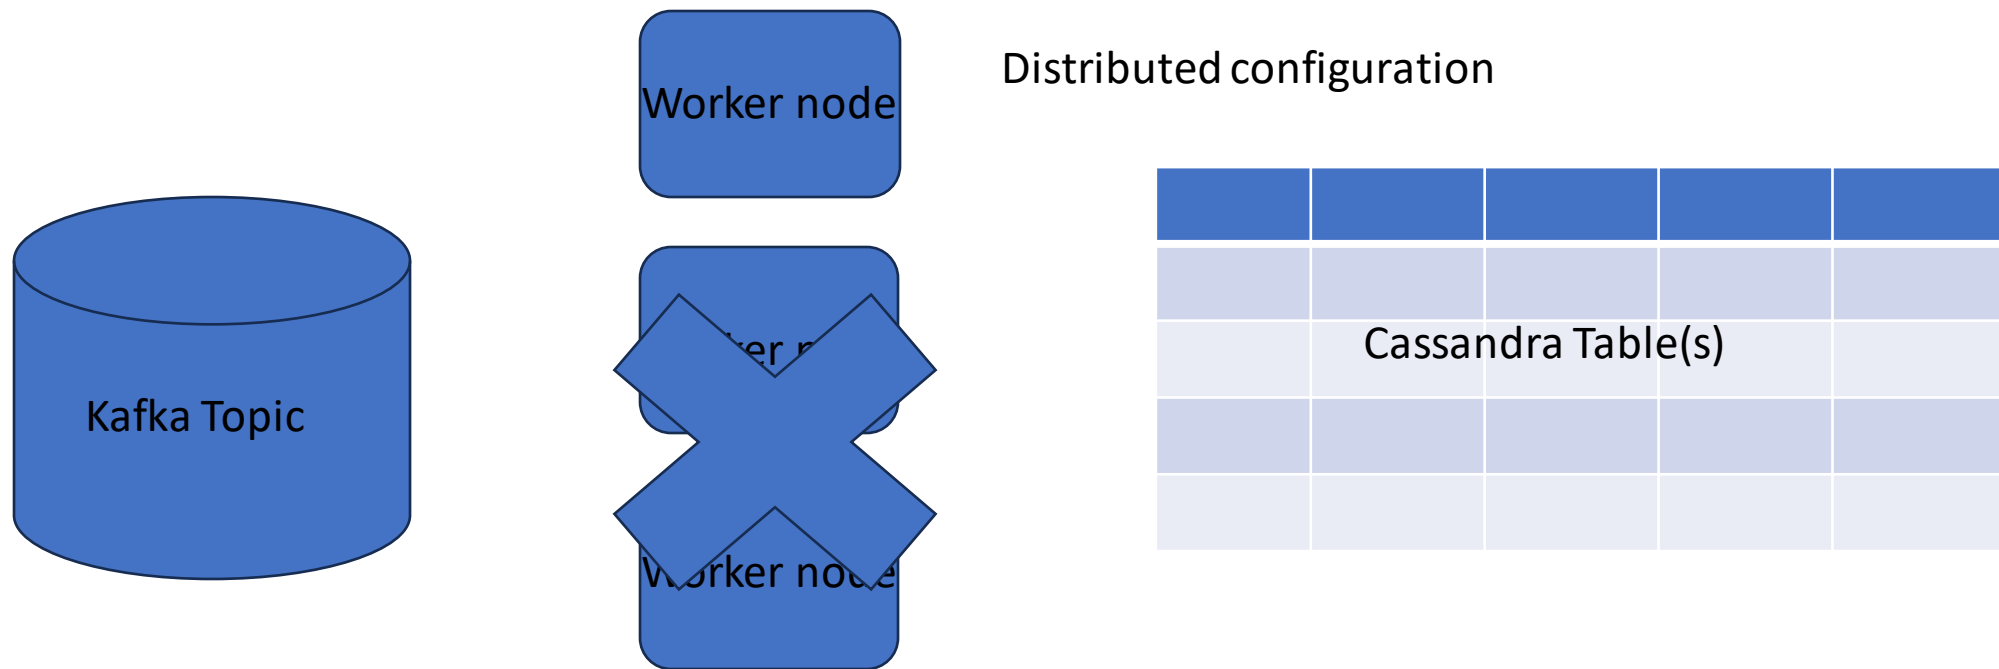
We have shown Cassandra as a Sink

- provision for 3 worker nodes running in distributed mode
- not much change in the various properties file
- just worry about the Kafka topics to Cassandra tables, that's it
- what about Cassandra as a source
- what is in bold face is what needs to change for Cassandra to be the source


Standalone configuration



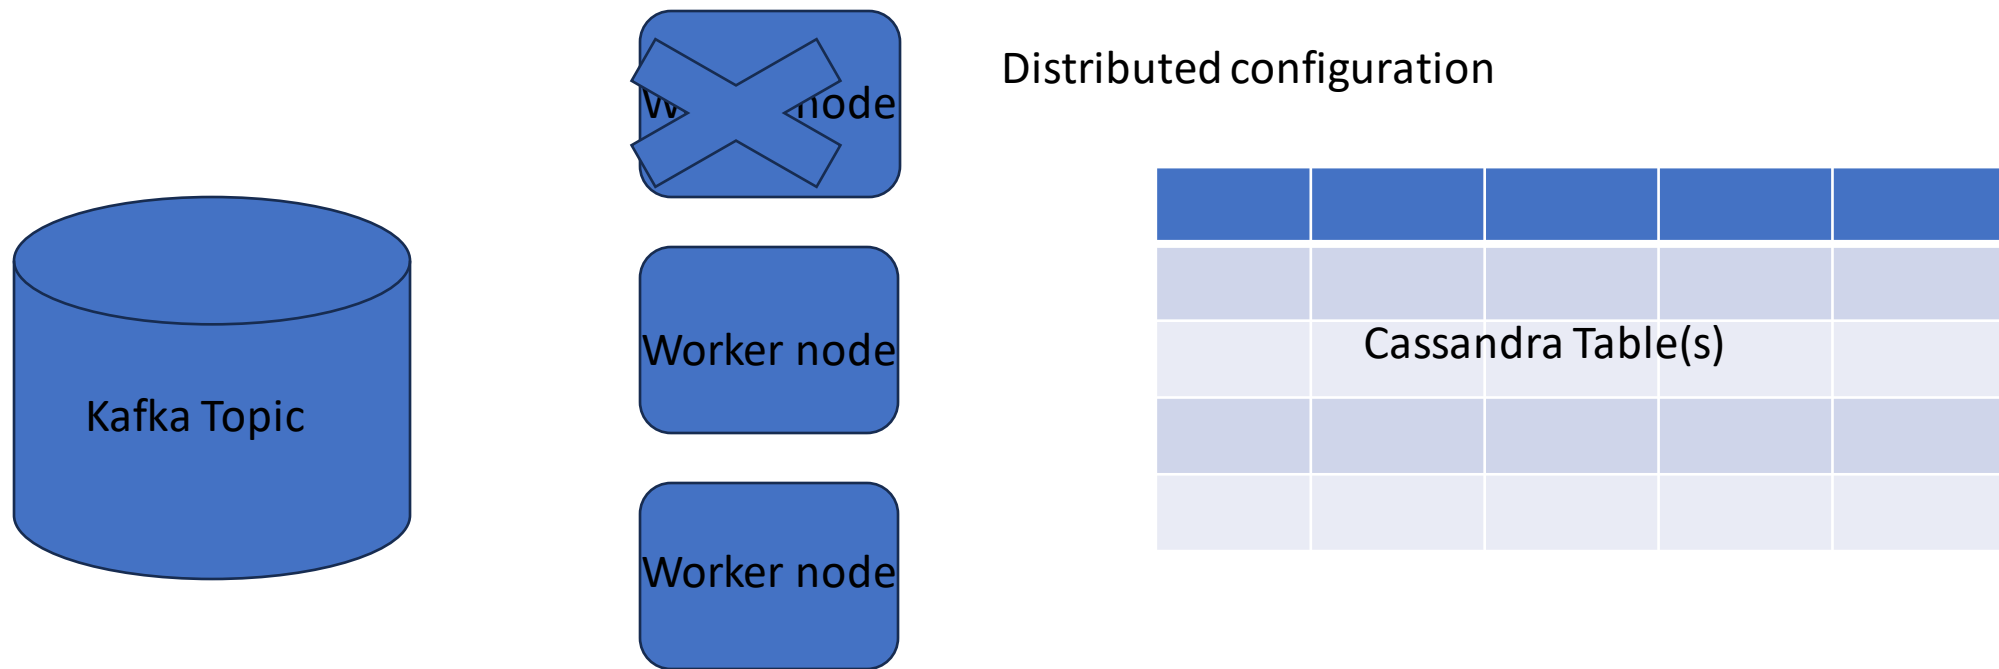
Distributed configuration




Standalone configuration



Distributed configuration



- Confluent has no official Cassandra Source connector available
- Check with github or open source for possible downloads
- BYO a lot of work; deep knowledge of Cassandra and Kafka internals to come to fruition
- KIP – Kafka Improvement Process it is like an RFP for the Kafka Developer Community
- For now, only the Cassandra Sink Connector is available



Thank you

The Brillio Team